ECE408 / CS483 / CSE408
Summer 2024

Applied Parallel Programming

Lecture 13: Atomic Operations
and Histogramming

What Will You Learn Today?

- atomic operations
  - read-modify-write in parallel computation
  - primitive form of "critical regions" in parallel programs
  - how to use in CUDA
  - why atomic operations reduce memory system throughput
  - how to avoid atomic operations in some parallel algorithms
- practical histogram programming techniques
  - basic histogram algorithm using atomic operations
  - atomic operation throughput
  - privatization

A Common Collaboration Pattern

- Multiple bank tellers count the total amount of cash in the safe
- Each grab a pile and count
- Have a central display of the running total
- Whenever someone finishes counting a pile, add the subtotal of the pile to the running total
- A bad outcome
  - Some of the piles were not accounted for

A Common Arbitration Pattern

- Multiple customers booking air tickets
- Each
  - Brings up a flight seat map
  - Decides on a seat
  - Update the the seat map, mark the seat as taken

- A bad outcome
  - Multiple passengers ended up booking the same seat

Read-Modify-Write Operations

thread1: Old ← Mem[x]   thread2: Old ← Mem[x]
       New ← Old + 1   New ← Old + 1
      Mem[x] ← New   Mem[x] ← New

If Mem[x] was **initially 0**, what would the value of Mem[x] be after threads 1 and 2 have completed?

– What does each thread get in their Old variable?

The answer may vary due to data races. To avoid data races, you should use atomic operations

Timing Scenario #1

Time	Thread 1	Thread 2
1	(0) Old ← Mem[x]	
2	(1) New ← Old + 1	
3	(1) Mem[x] ← New	
4		(1) Old ← Mem[x]
5		(2) New ← Old + 1
6		(2) Mem[x] ← New

- Thread 1 Old = 0
- Thread 2 Old = 1
- Mem[x] = 2 after the sequence

Timing Scenario #2

Time	Thread 1	Thread 2
1		(0) Old ← Mem[x]
2		(1) New ← Old + 1
3		(1) Mem[x] ← New
4	(1) Old ← Mem[x]	
5	(2) New ← Old + 1	
6	(2) Mem[x] ← New	

- Thread 1 Old = 1
- Thread 2 Old = 0
- Mem[x] = 2 after the sequence

Timing Scenario #3

Time	Thread 1	Thread 2
1	(0) Old ← Mem[x]	
2	(1) New ← Old + 1	
3		(0) Old ← Mem[x]
4	(1) Mem[x] ← New	
5		(1) New ← Old + 1
6		(1) Mem[x] ← New

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

Timing Scenario #4

Time	Thread 1	Thread 2
1		(0) Old ← Mem[x]
2		(1) New ← Old + 1
3	(0) Old ← Mem[x]	
4		(1) Mem[x] ← New
5	(1) New ← Old + 1	
6	(1) Mem[x] ← New	

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

Atomic Operations Prevent Interleaving

thread1:Old ← Mem[x]
New ← Old + 1
Mem[x] ← New

thread2:Old ← Mem[x]
New ← Old + 1
Mem[x] ← New

Or

thread2:Old ← Mem[x]
New ← Old + 1
Mem[x] ← New

thread1:Old ← Mem[x]
New ← Old + 1
Mem[x] ← New

Without Atomic Operations

Mem[x] initialized to 0

thread1: Old ← Mem[x]

                                             thread2: Old ← Mem[x]

    New ← Old + 1

    Mem[x] ← New   New ← Old + 1

                                     Mem[x] ← New

- Both threads receive 0
- Mem[x] becomes 1

Needed When Threads Write to Same Location

When **two threads**

- **may write to** the **same memory** location,

- the program may **need atomic operations**.

Sharing is **not always easy to recognize**…

- Do two insertions into a hash table share data?

- What about two graph node updates
based on all of the nodes' neighbors?

- What if nodes are on same side of bipartite graph?

What Exactly is "Atomic?"

To a high-energy photon, atoms are not.

Atomicity is ALWAYS with respect to something.

Two **sections of code**

- **that execute atomically with respect to one another**
- **appear** to the software **as though**
- the programs' **execution did not interleave** at all.

What Can Go Wrong?

Common failure mode:

- **Programmer *thinks* operations** are **independent**.
- Hasn't considered input data for which they are not.
- Or another programmer reuses code without understanding assumptions that imply independence.

Also: **atomicity does not constrain relative order**.

Implementing Atomic Operations

- Many ISAs offer synchronization primitives,
  - **instructions with** one (or more) **address operands**
  - **that execute atomically with respect to one another** when used on the same address.
- Mostly read, modify, write operations
  - Bit test and set
  - Compare and swap / exchange
  - Swap / exchange
  - Fetch and increment / add

Atomicity Enforced by Microarchitecture

When synchronization primitives execute,

- **hardware ensures** that **no other** thread
- **accesses** the location **until** the operation is **complete**.

Other threads that access the location

- are typically stalled or held in a queue until their turn.
- **Threads perform atomic operations serially**.

Atomic Operations in CUDA

- Function calls that are translated into single ISA instructions (a.k.a. *intrinsics*)
  - Atomic add, sub, inc, dec, min, max, exch (exchange), CAS (compare and swap)
  - Read CUDA C programming Guide for more details

- Atomic Add

    int atomicAdd(int **address**, int **val**);*

  reads the 32-bit word **old** pointed to by **address** in global or shared memory, computes **(old + val)**, and stores the result back to memory at the same address. The function returns **old**.

More Atomic Adds in CUDA

- Unsigned 32-bit integer atomic add

  unsigned int atomicAdd(unsigned int address, unsigned int val);*

- Unsigned 64-bit integer atomic add

  unsigned long long int atomicAdd(unsigned long long int address, unsigned long long int val);*

- Single-precision floating-point atomic add (capability > 2.0)
  - float atomicAdd(float* address, float val);

Building Synchronization with Atomics

- How would we build `__syncthreads()` for a thread block?
- How would we create `__syncthreads()` for an entire grid?
  - And why would this not be a good idea?

- How would we create a critical section? That is, one thread per block executing a particular section of code?

- How would we create a critical section per grid?
  - Why doesn't this have the same issue as `__syncthreads()` for grid?

Atomic Compare and Swap (CAS)

```
Bool atomicCAS(int *address, int old, int new)
{
    if (*address != old)
        return false;

    *address = new;

    return true;
}
```

Atomic Add using CAS

```
int atomicAdd(int* address, int value)
{
  Bool done = false;

  while (!done) {
    old_v = *address;
    done = atomicCAS(address, old_v, old_v+value);
  }

  return old_v;
}
```

Histogramming

- A method for extracting notable features and patterns from large data sets
  - Feature extraction for object recognition in images
  - Fraud detection in credit card transactions
  - Correlating heavenly object movements in astrophysics
  - …

- Basic histograms - for each element in the data set, use the value to identify a "bin" to increment

A Histogram Example

- In sentence "Programming Massively Parallel Processors" build a histogram of frequencies of each letter

- A(4), C(1), E(1), G(1), …

- How do you do this in parallel?
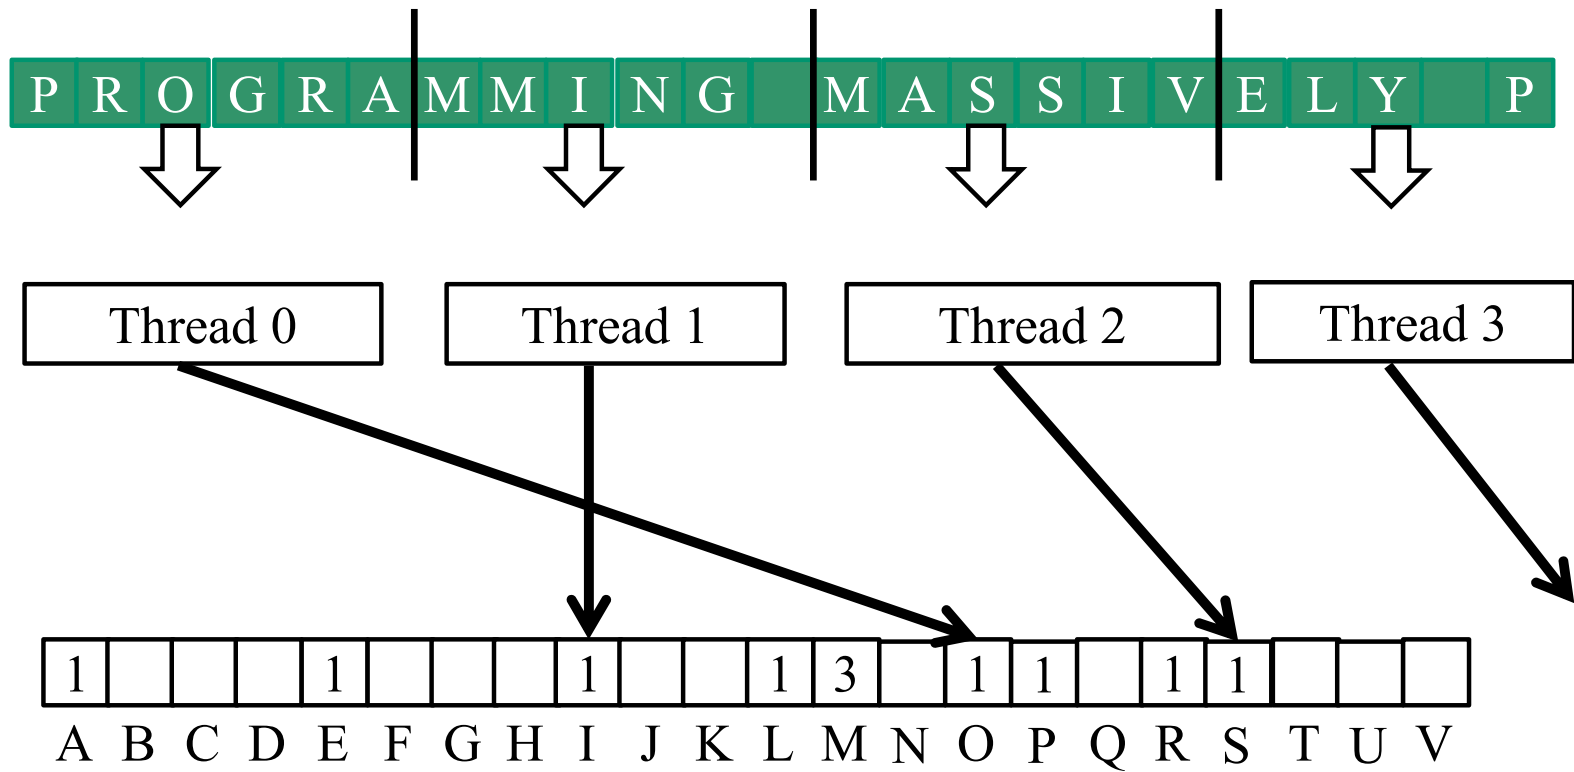
Iteration #1 – 1st letter in each section
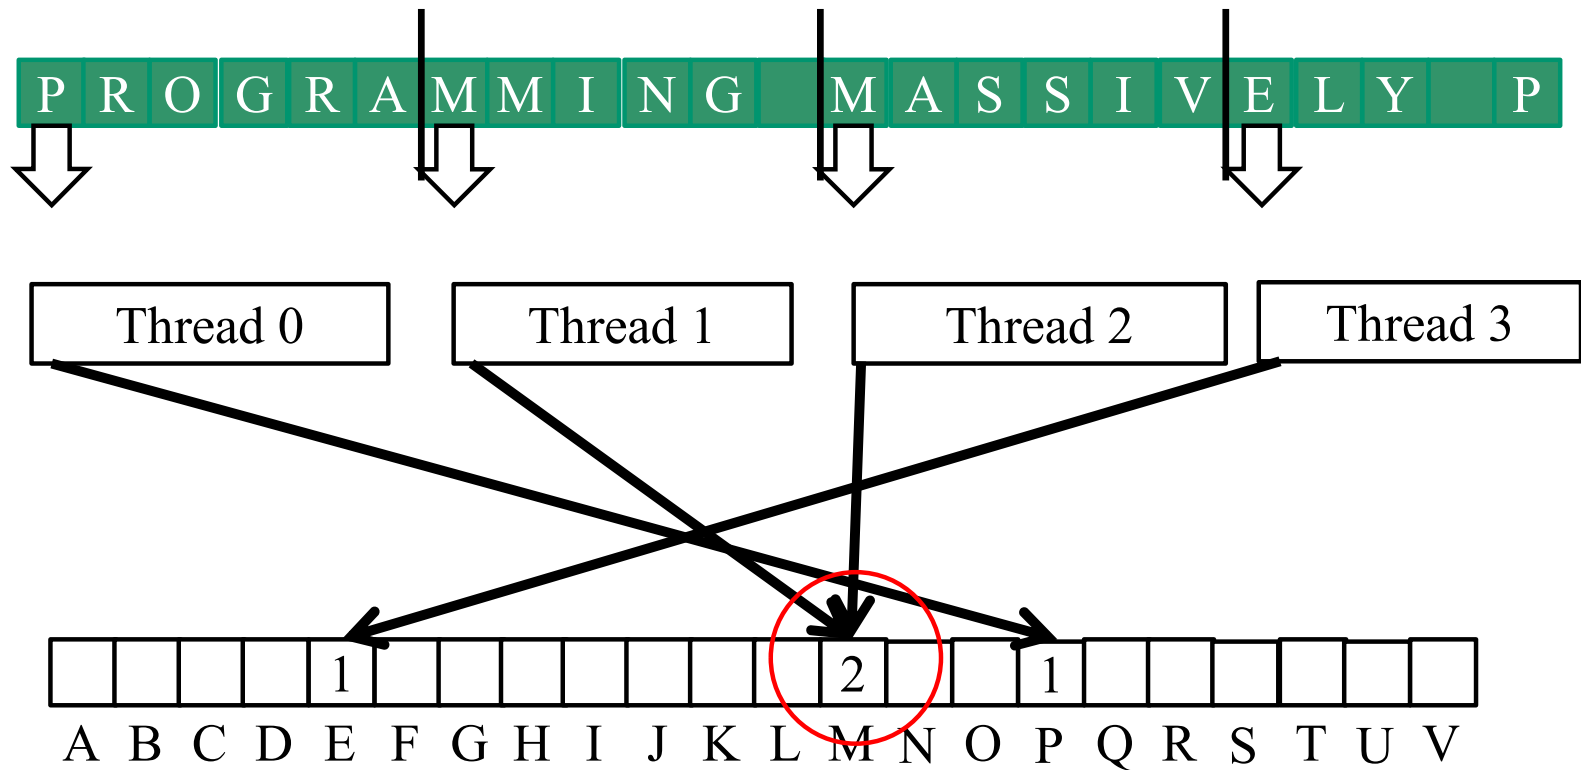
Atomic operation enforces correct update

Iteration #2 – 2nd letter in each section

Iteration #3

Iteration #4
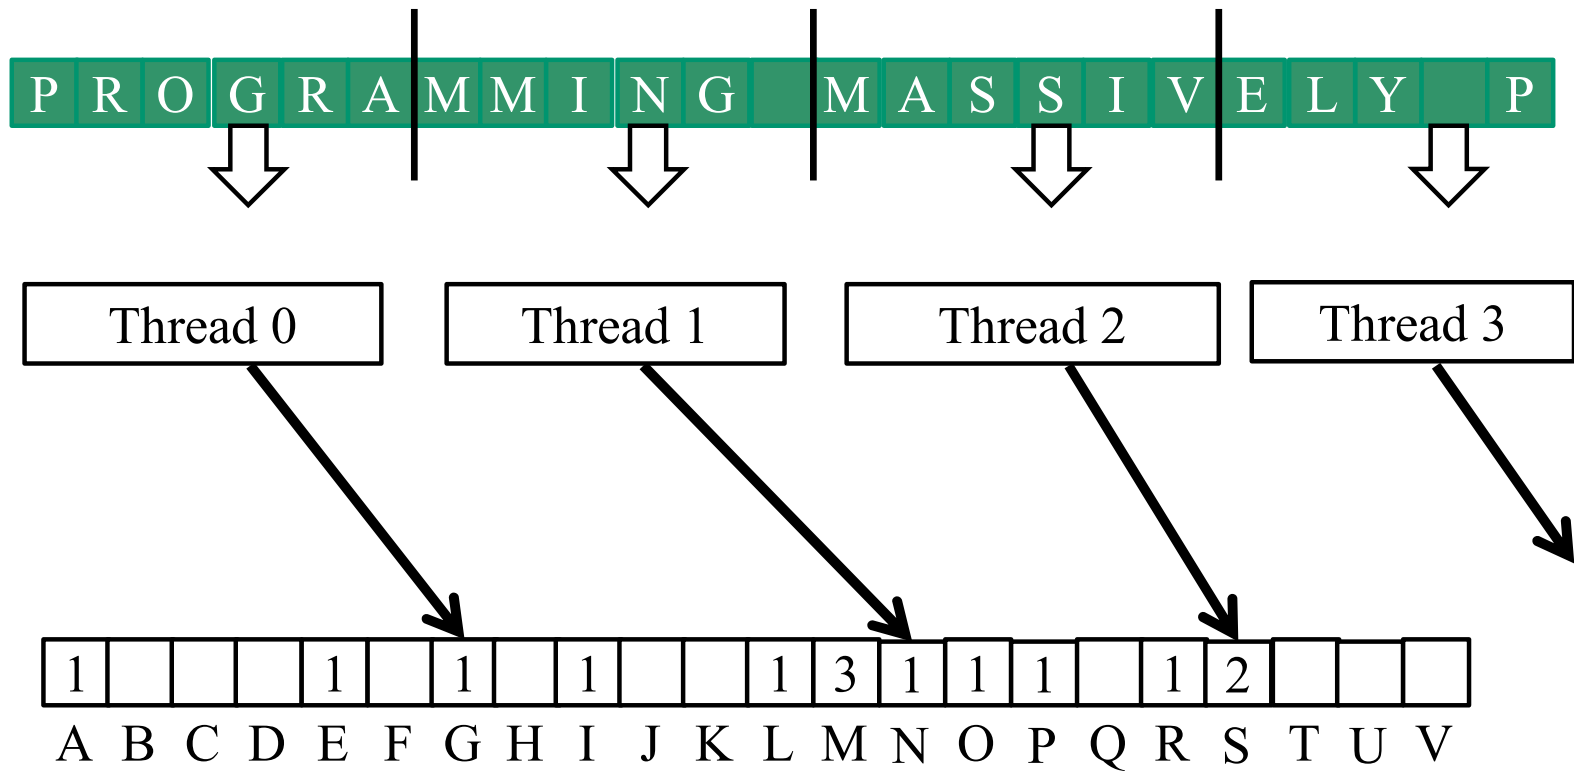
Iteration #5

What is wrong with the algorithm?

What is wrong with the algorithm?
A better approach:

- Reads from the input array are not coalesced
    - Assign inputs to each thread in a strided pattern
    - Adjacent threads process adjacent input letters

© David Kirk/N
of Illinois, 2007-
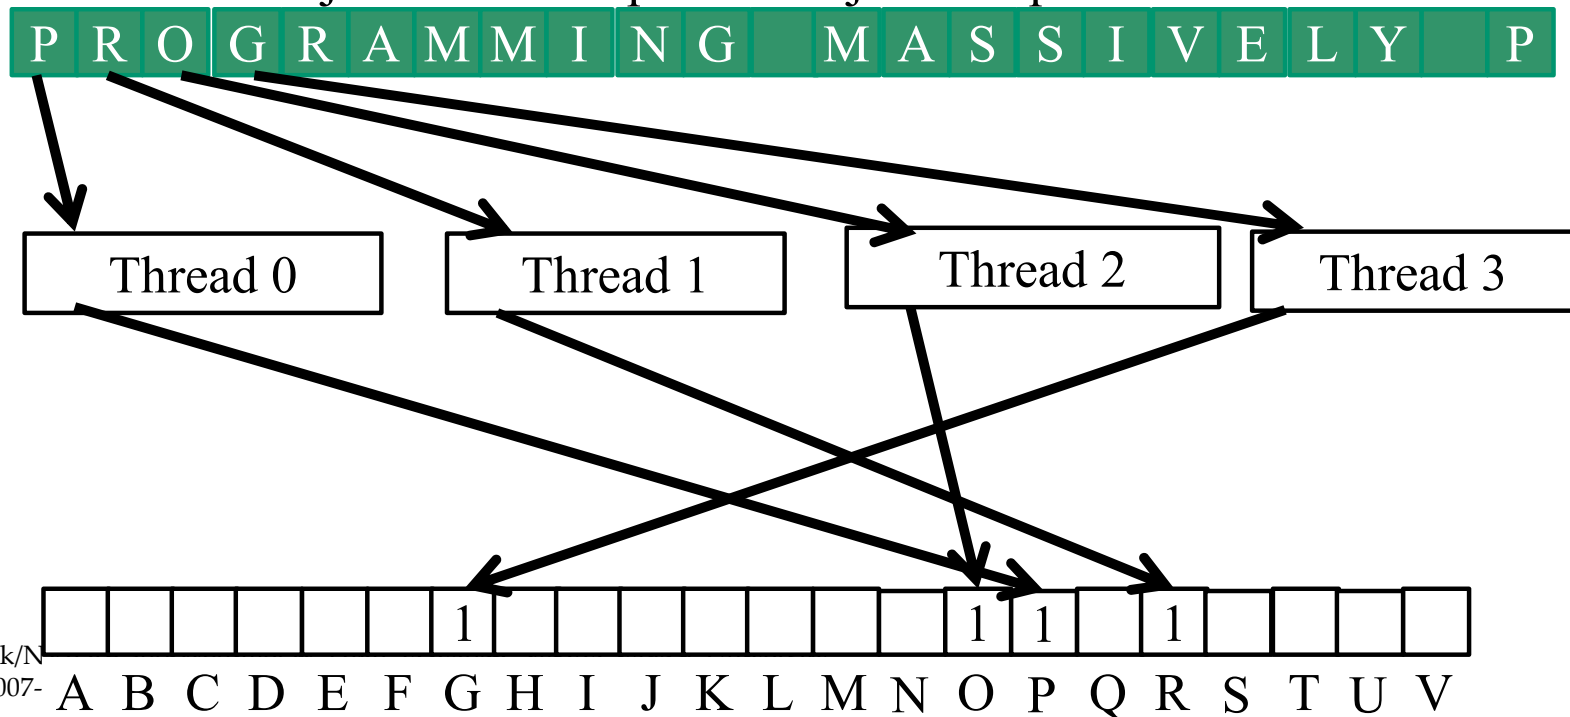
Iteration 2

- All threads move to the next section of input
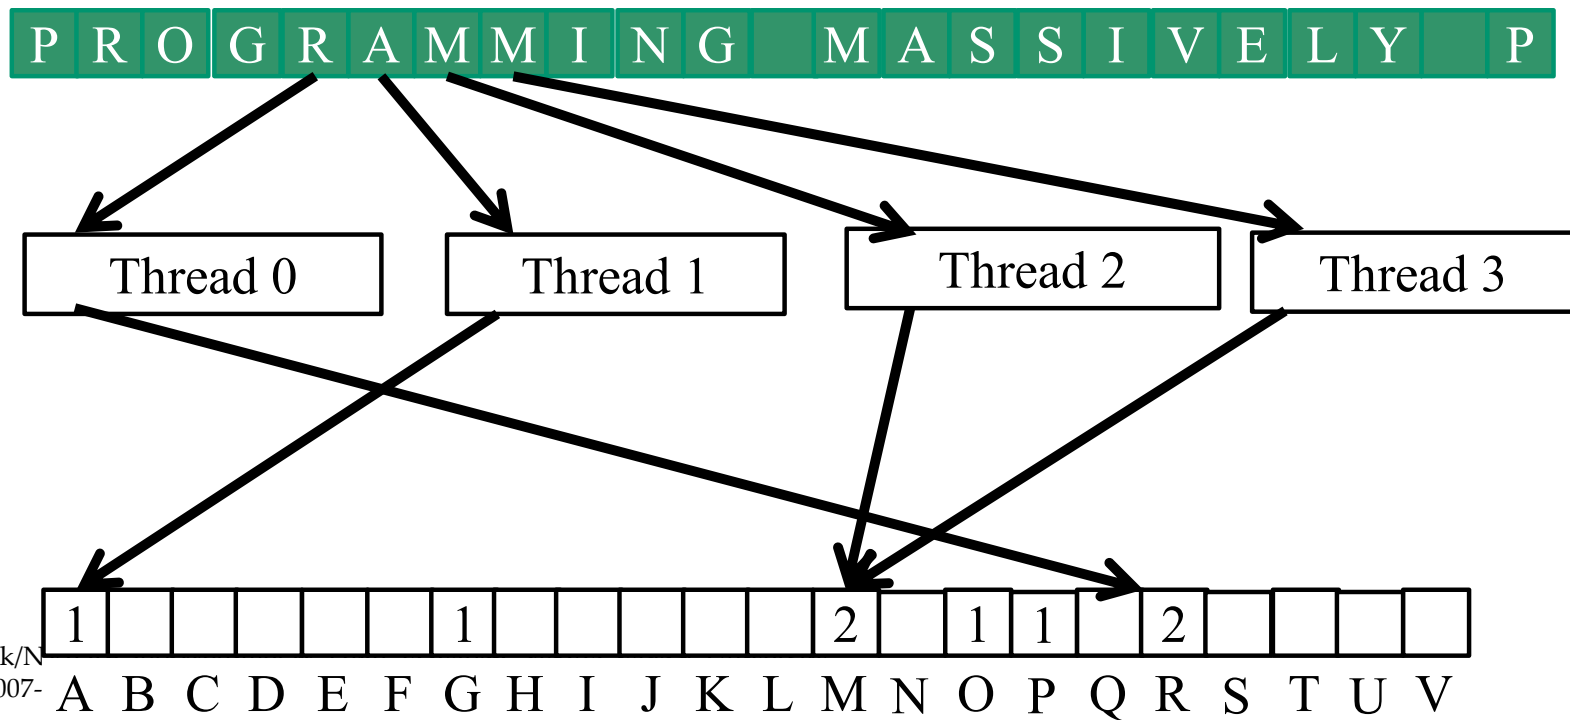
A Histogram Kernel

- The kernel receives a pointer to the input buffer
- Each thread process the input in a strided pattern

```
__global__ void histo_kernel(unsigned char *buffer,
                             long size, unsigned int *histo)
{
    int i = threadIdx.x + blockIdx.x * blockDim.x;

    // stride is total number of threads
    int stride = blockDim.x * gridDim.x;
```

More on the Histogram Kernel

```
// All threads in the grid collectively handle
// blockDim.x * gridDim.x consecutive elements


while (i < size) {
    atomicAdd( &(histo[buffer[i]]), 1);
    i += stride;
  }
}
```

Atomic Operations on DRAM

DRAM Bank	-	DRAM Channel	Intra-channel	Burst	Byte Offset
A[21:20]	A[19:14]	A[13:11]	A[10:6]	A[5:2]	A[1:0]

SM#0 SM#1 ...... SM#29

Interconnect

DRAM Ctrlr 0 DRAM Ctrlr 1 DRAM Ctrlr 2 DRAM Ctrlr 3 ..... DRAM Ctrlr 7

Bank Selection

DRAM Bank X 32

32 Interleaved ports per DRAM controller, selected by A[10:6]

4 Banks per interleaved port, selected by A[21:20]
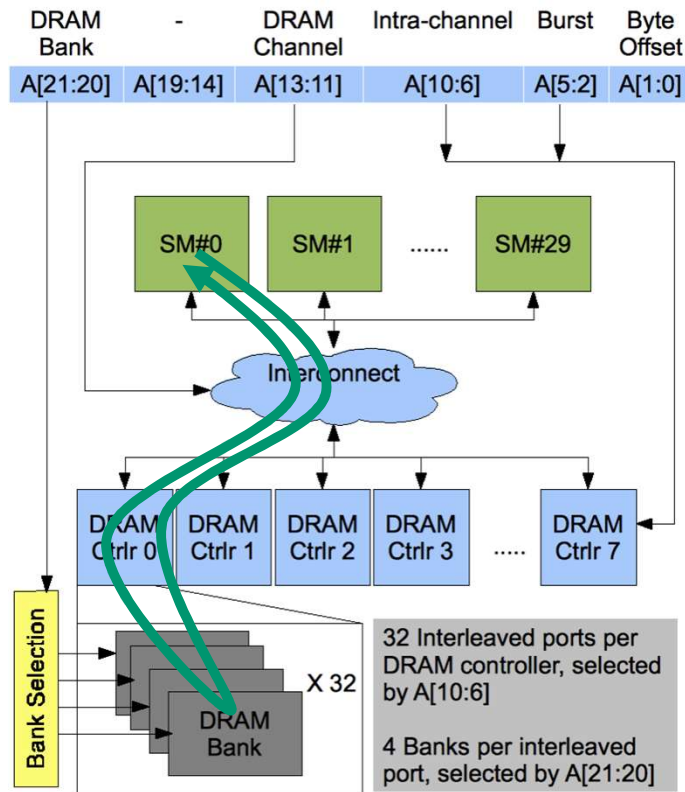
- An atomic operation starts with a read, with a latency of a few hundred cycles

Atomic Operations on DRAM

DRAM Bank	-	DRAM Channel	Intra-channel	Burst	Byte Offset
A[21:20]	A[19:14]	A[13:11]	A[10:6]	A[5:2]	A[1:0]

SM#0   SM#1   ......   SM#29

Interconnect

DRAM Ctrlr 0   DRAM Ctrlr 1   DRAM Ctrlr 2   DRAM Ctrlr 3   .....   DRAM Ctrlr 7

Bank Selection

X 32   DRAM Bank

32 Interleaved ports per DRAM controller, selected by A[10:6]

4 Banks per interleaved port, selected by A[21:20]
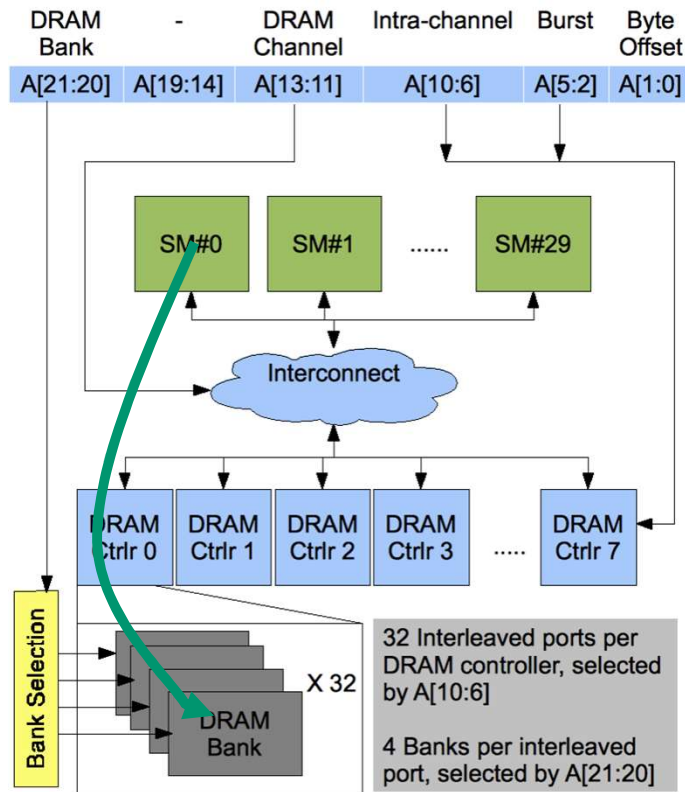
- An atomic operation starts with a read, with a latency of a few hundred cycles
- The atomic operation ends with a write, with a latency of a few hundred cycles
- During this whole time, no one else can access the location

Atomic Operations on DRAM

- Each Load-Modify-Store has two full memory access delays
  - All atomic operations on the same variable (RAM location) are serialized
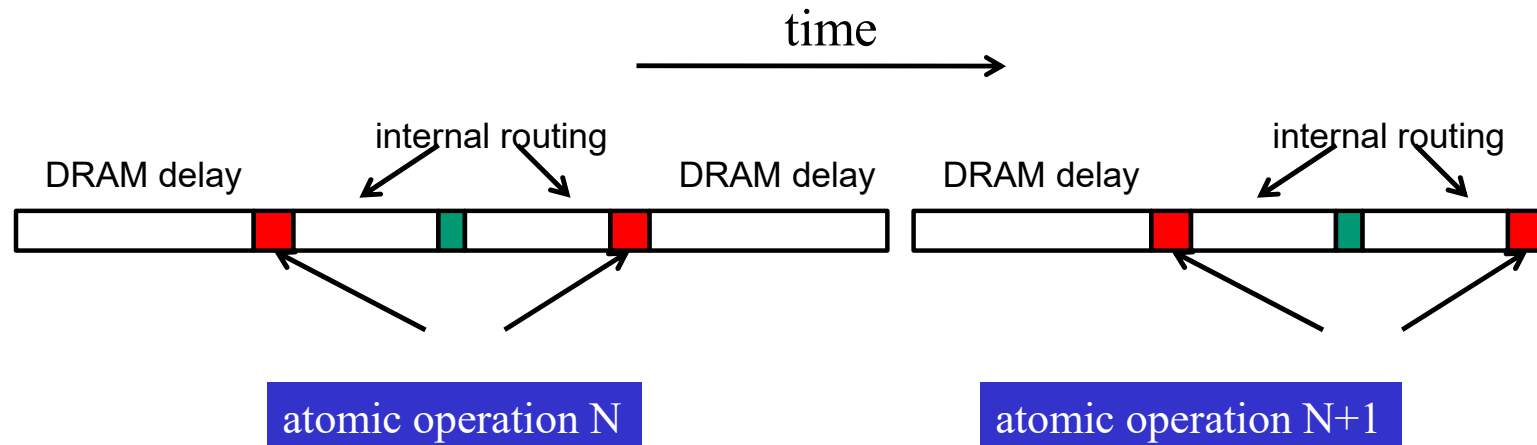
Latency determines throughput of atomic operations

- Throughput of an atomic operation is the rate at which the application can execute an atomic operation on a particular location.

- The rate is limited by the total latency of the read-modify-write sequence, typically more than 1000 cycles for global memory (DRAM) locations.

- This means that if many threads attempt to do atomic operation on the same location (contention), the memory bandwidth is reduced to < 1/1000!

Waiting for Atomics, or Waiting in a Queue?

Imagine yourself in a supermarket with a single cashier…

- The person in front of your forgot an item
- They run to get the item while the line waits
- The rate of people checking out is reduced due to the long latency of running to the isle and back.

Now imagine a store in which every customer starts checking out before fetching any of their items…

The rate of the checkout will be 1 / (entire shopping time of each customer)

Hardware Improvements

- Atomic operations on L2 cache
  - medium latency, but still serialized
  - Global to all blocks
  - "Free improvement" on Global Memory atomics

time

internal routing

atomic operation N   atomic operation N+1

40

Hardware Improvements

- Atomic operations on Shared Memory
  - Very short latency, but still serialized
  - Private to each thread block
  - Need algorithm work by programmers (more later)

time

internal routing

atomic operation N   atomic operation N+1

41

Atomics in Shared Memory Requires Privatization

- Create private copies of the histo[] array for each thread block

__global__ void histo_kernel(unsigned char *buffer,

long size, unsigned int *histo)

{

    __shared__ unsigned int histo_private[256];

  if (threadIdx.x < 256) histo_private[threadidx.x] = 0;

    __syncthreads();

Build Private Histogram

```
    int i = threadIdx.x + blockIdx.x * blockDim.x;
// stride is total number of threads
    int stride = blockDim.x * gridDim.x;
    while (i < size) {
        atomicAdd( &(private_histo[buffer[i]), 1);
        i += stride;
    }
```

Build Final Histogram

```
// wait for all other threads in the block to finish
__syncthreads();

if (threadIdx.x < 256)
  atomicAdd( &(histo[threadIdx.x]),
                          private_histo[threadIdx.x] );

}
```

More on Privatization

- Privatization is a powerful and frequently used techniques for parallelizing applications

- The operation needs to be associative and commutative
  - Histogram add operation is associative and commutative

- The histogram size needs to be small
  - Fits into shared memory

- What if the histogram is too large to privatize?

QUESTIONS?

READ CHAPTER 9!

Problem Solving

- Q: Suppose a processor supports atomic operations in L2 cache. Assume that each atomic operation takes 5ns to complete in L2 cache and 120ns to complete in DRAM. The kernel performs 20 floating-point operations per atomic operation, and a floating-point operation takes 1ns. **Assume the time of L2 atomic operations, DRAM atomic operations, and floating-point operations in each thread do not overlap.*** The floating-point throughput of the kernel execution is 0.2424 GFLOPS, and every thread in a block performs 5 atomic operations and 100 floating-point operations. What percent of the atomic operations happened in L2 cache?

* An unusual assumption on a GPU, but simplifies the question!

- A: 50%

Problem Solving - Explanation

- The total time taken by a thread for all operations, assuming all atomic operations are in L2 cache, is: 5 atomic operations * 5ns/atomic operation + 100 floating-point operations * 1ns/floating-point operation = 25ns + 100ns = 125ns

- Similarly, the total time taken by a thread for all operations, assuming all atomic operations are in DRAM, is: 5 atomic operations * 120ns/atomic operation + 100 floating-point operations * 1ns/floating-point operation = 600ns + 100ns = 700ns

- Let's denote the percentage of atomic operations that happened in DRAM as x. Therefore, the percentage of atomic operations that happened in L2 cache is $1 - x$ and the total execution time of a thread is: $(1 - x)$ * 125ns/thread + x * 700ns/thread

- Given that the floating-point throughput of the kernel execution is 0.2424 GFLOPS, we can calculate the total time taken for all operations in a thread (1 GFLOP = 1 billion floating-point operations): 0.2424 GFLOPS = 0.2424 * 1 billion floating-point operations/second = 242.4 million floating-point operations/second.

- Since every thread performs 100 floating-point operations, the number of threads that can be executed per second is: 242.4 million threads/second = 242.4 million floating-point operations/second / 100 floating-point operations/thread = 2.424 million threads/second.

- Therefore, the total time taken for all operations in a thread is: 1 / 2.424 million threads/second = 0.0000004125 seconds/thread = 412.5ns/thread.

- Thus: $(1 - x)$ * 125ns/thread + x * 700ns/thread = 412.5ns/thread. Solving this for x gives us 50%.