Two vertical lines, one blue and one orange, are positioned on the left side of the slide.

ECE408 / CS483 / CSE408  
Summer 2024


Applied Parallel Programming

Lecture 18: Computation in Deep  
Neural Networks


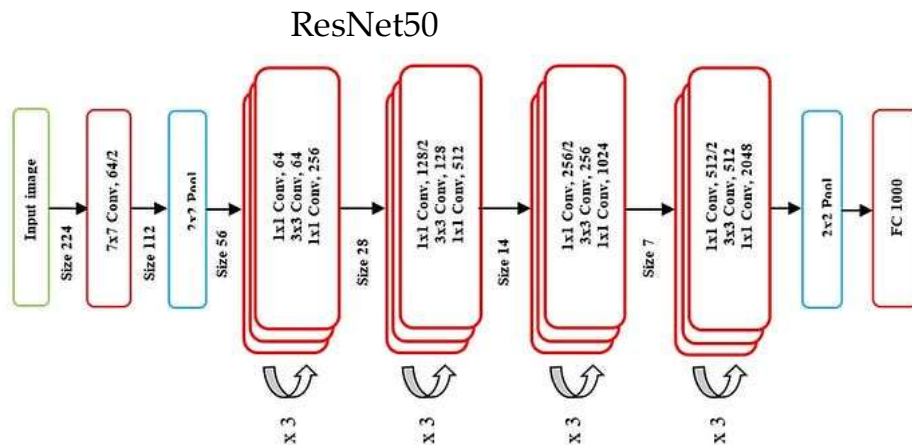
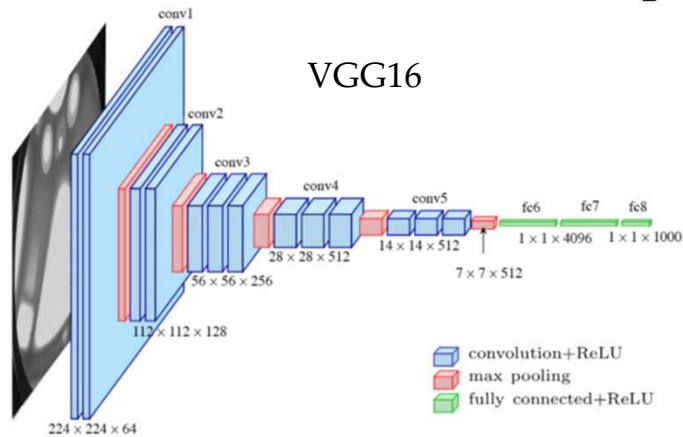
# What Will You Learn Today?

to implement the different types of layers  
in a Convolutional Neural Network (CNN)

# LeNet-5: CNN for hand-written digit recognition



# Many Types of CNNs



# Anatomy of a Convolution Layer

Input features


- A inputs each  $N_1 \times N_2$

Convolution Layer

- B convolution kernels each  $K_1 \times K_2$

Output Features (total of B)

- $A \times B$  outputs each  
 $(N_1 - K_1 + 1) \times (N_2 - K_2 + 1)$



# Notion of a Channel in Input Layer

Some Set of Input Features are Related


- For example: Red, Green, Blue

Convolution Layer

- Different kernels per channel


Output Features

- Channels combine per output feature




# 2-D Pooling (Subsampling)

- A subsampling layer
  - Sometimes with bias and non-linearity built in
- Common types
  - max, average,  $L^2$  norm, weighted average
- Helps make representation invariant to size scaling and small translations in the input

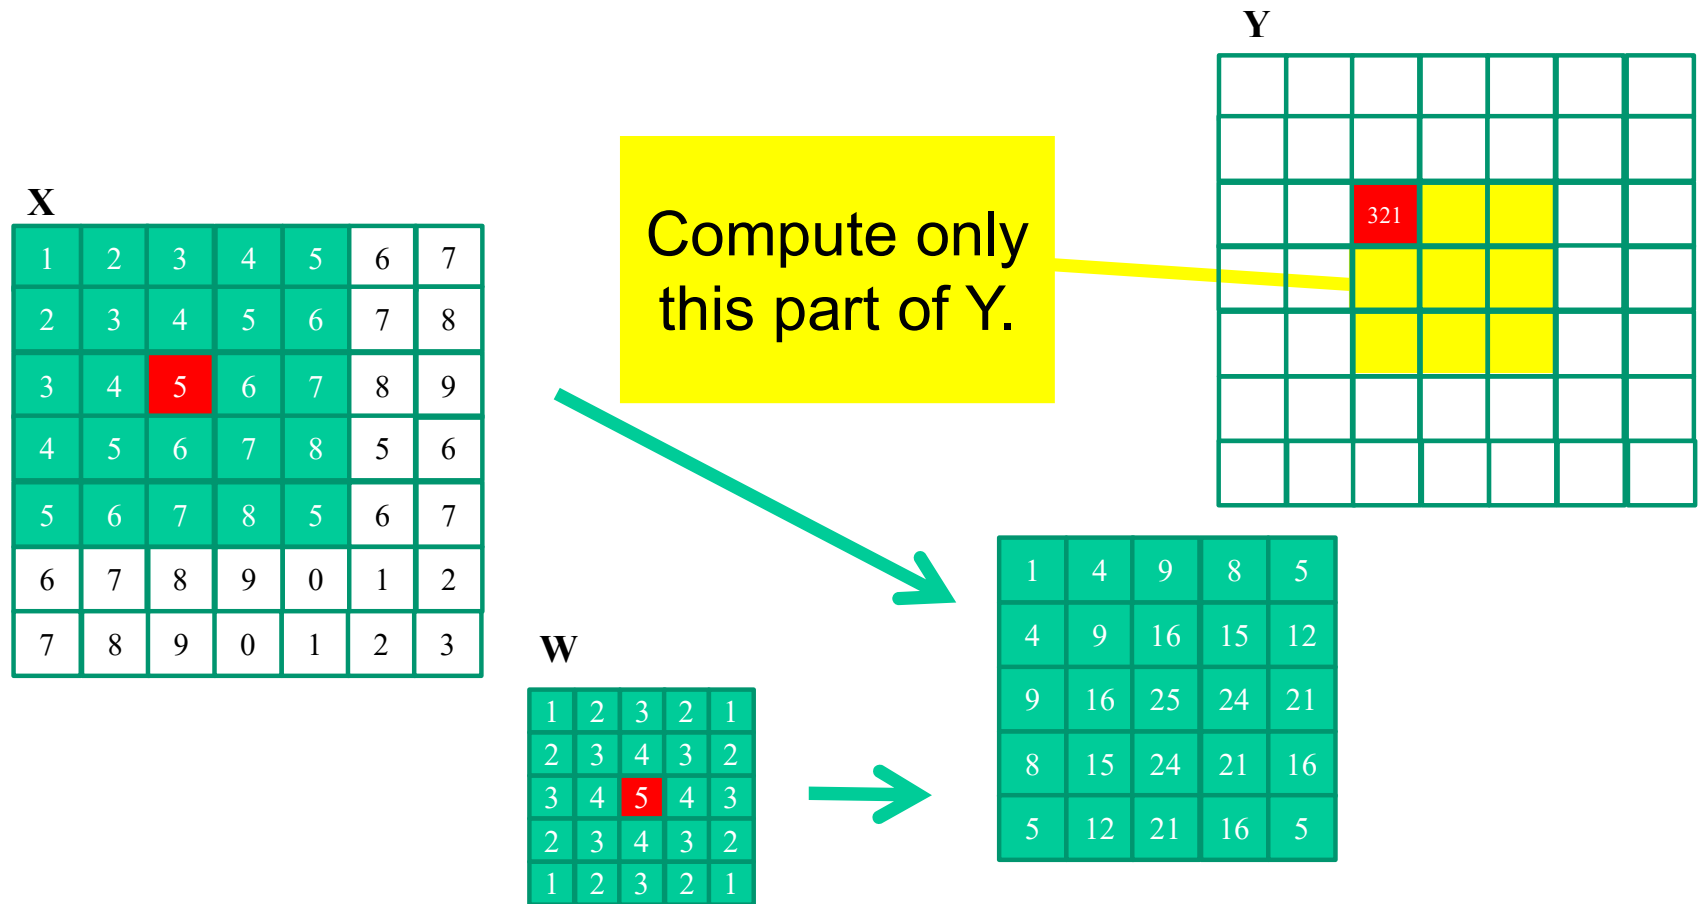


# Forward Propagation





# Outputs Must Use Full Mask/Kernel



# Example of the Forward Path of a Convolution Layer

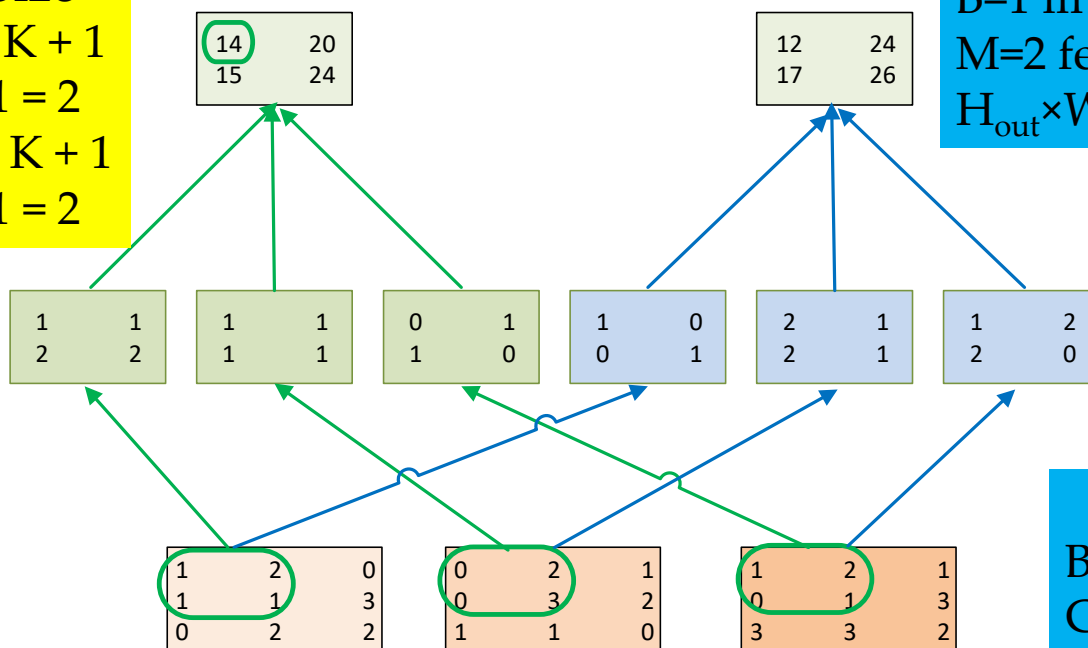
## Output Size

$$H_{\text{out}} = H - K + 1$$

$$= 3 - 2 + 1 = 2$$

$$W_{\text{out}} = W - K + 1$$

$$= 3 - 2 + 1 = 2$$



## Convolution Output Y

B=1 image

M=2 features per image

$H_{\text{out}} \times W_{\text{out}} = 2 \times 2$  values per feature

## Weights W

M=2 feature maps

C=3 channels per map

$K \times K = 2 \times 2$  pixels per channel

## Input X

B=1 image

C=3 channels

$H \times W = 3 \times 3$  pixels per channel

# Sequential Code: Forward Convolutional Layer

```
void convLayer_forward(int B, int M, int C, int H, int W, int K, float* X, float* W, float* Y) {
    int H_out = H - K + 1;           // calculate H_out, W_out
    int W_out = W - K + 1;

    for (int b = 0; b < B; ++b)      // for each image
        for(int m = 0; m < M; m++)   // for each output feature map
            for(int h = 0; h < H_out; h++) // for each output value (two loops)
                for(int w = 0; w < W_out; w++) {
                    Y[b, m, h, w] = 0.0f; // initialize sum to 0
                    for(int c = 0; c < C; c++) // sum over all input channels
                        for(int p = 0; p < K; p++) // KxK filter
                            for(int q = 0; q < K; q++)
                                Y[b, m, h, w] += X[b, c, h + p, w + q] * W[m, c, p, q];
                }
    }
```

# A Small Convolution Layer Example

Image  $b$  in mini batch

$x[b,0,_,_]$

1	2	0	1
1	1	3	2
0	2	2	0
2	1	0	3

1	1	1
2	2	3
2	1	0

$w[0,0,_,_]$

$x[b,1,_,_]$

0	2	1	0
0	3	2	1
1	1	0	2
2	1	0	3

1	2	3
1	1	0
3	0	1

$w[0,1,_,_]$

0	?
?	?

$y[b,0,_,_]$

$x[b,2,_,_]$

1	2	1	0
0	1	3	2
3	3	2	0
1	3	2	0

0	1	1
1	0	2
1	2	1

$w[0,2,_,_]$

$X[b, 1,_,_]$


$W[0,1,_,_]$

$Y[b, 0,_,_]$

output map


# A Small Convolution Layer Example

$c = 0$




# A Small Convolution Layer Example

$c = 1$



# A Small Convolution Layer Example

$c = 2$



# Parallelism in a Convolution Layer

**Output feature maps** can be calculated in parallel

- Usually a small number, not sufficient to fully utilize a GPU

All **output** feature map **pixels** can be calculated in parallel

- All rows can be done in parallel
- All pixels in each row can be done in parallel
- Large number but diminishes as we go into deeper layers

All **input feature maps** can be processed in parallel,  
but need atomic operation or tree reduction (we'll learn later)

**Different layers may demand different strategies.**



# Subsampling (Pooling) by Scale N

## Convolution Output Y

B images

M features per image

$H_{\text{out}} \times W_{\text{out}}$  values per feature

Average over  $N \times N$   
blocks,

then calculate sigmoid

## Subsampling/Pooling Output S

B images

M features per image

$H_{S(N)} \times W_{S(N)}$  values per feature

## Output Size

$$H_{S(N)} = \text{floor}(H_{\text{out}} / N)$$

$$W_{S(N)} = \text{floor}(W_{\text{out}} / N)$$

# Sequential Code: Forward Pooling Layer

```
void poolingLayer_forward(int B, int M, int H_out, int W_out, int N, float* Y, float* S)
{
    for (int b = 0; b < B; ++b)                // for each image
        for (int m = 0; m < M; ++m)            // for each output feature map
            for (int x = 0; x < H_out/N; ++x)    // for each output value (two loops)
                for (int y = 0; y < W_out/N; ++y) {
                    float acc = 0.0f                // initialize sum to 0
                    for (int p = 0; p < N; ++p)    // loop over NxN block of Y (two loops)
                        for (int q = 0; q < N; ++q)
                            acc += Y[b, m, N*x + p, N*y + q];
                    acc /= N * N;                // calculate average over block
                    S[b, m, x, y] = sigmoid(acc + bias[m]) // bias, non-linearity
                }
    }
```


# Kernel Implementation of Subsampling Layer

- Straightforward mapping from grid to subsampled output feature map pixels
- in GPU kernel,
  - need to manipulate index mapping
  - for accessing the output feature map pixels
  - of the previous convolution layer.
- Often merged into the previous convolution layer to save memory bandwidth

# Design of a Basic Kernel

- Each block computes
  - a tile of output pixels for one feature
  - TILE\_WIDTH pixels in each dimension
- Grid's X dimension maps to M output feature maps
- Grid's Y dimension maps to the tiles in the output feature maps (linearized order).
- (Grid's Z dimension is used for images in batch, which we omit from slides.)

tiles covering an  
output feature map,  
marked with  
linearized indices



0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19

# A Small Example


Assume

- **M = 4** (4 output feature maps),
- thus 4 blocks in the X dimension, and
- **W\_out = H\_out = 8** (8x8 output features).

If **TILE\_WIDTH = 4**,  
we also need 4 blocks in the Y dimension:

- for each output feature,
- top two blocks in each column calculates the top row of tiles, and
- bottom two calculate the bottom row.

CUDA Grid and  
Threadblocks



Output Feature Maps and Tiles <sup>22</sup>

# Host Code for a Basic Kernel: CUDA Grid

Consider an output feature map:

- width is **W\_out**, and
- height is **H\_out**.
- Assume these are multiples of **TILE\_WIDTH**.

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19

Let **X\_grid** be the number of blocks needed in X dim (5 above).

Let **Y\_grid** be the number of blocks needed in Y dim (4 above).

# Host Code for a Basic Kernel: CUDA Grid

(Assuming  $W_{out}$  and  $H_{out}$  are multiples of  $TILE\_WIDTH$ .)

```
#define TILE_WIDTH 16          // We will use 4 for small examples.  
W_grid = W_out/TILE_WIDTH;    // number of horizontal tiles per output map  
H_grid = H_out/TILE_WIDTH;    // number of vertical tiles per output map  
Y = H_grid * W_grid;  
  
dim3 blockDim(TILE_WIDTH, TILE_WIDTH, 1); // output tile for untiled code  
dim3 gridDim(M, Y, 1);  
  
ConvLayerForward_Kernel<<< gridDim, blockDim >>>(...);
```

# Sequential Code: Forward Convolutional Layer

```
void convLayer_forward(int B, int M, int C, int H, int W, int K, float* X, float* W, float* Y) {  
    int H_out = H - K + 1;           // calculate H_out, W_out  
    int W_out = W - K + 1;
```

```
    for (int b = 0; b < B; ++b)      // for each image  
        for(int m = 0; m < M; m++)    // for each output feature map  
            for(int h = 0; h < H_out; h++) // for each output value (two loops)  
                for(int w = 0; w < W_out; w++) {  
                    Y[b, m, h, w] = 0.0f; // initialize sum to 0  
                    for(int c = 0; c < C; c++) // sum over all input channels  
                        for(int p = 0; p < K; p++) // KxK filter  
                            for(int q = 0; q < K; q++)  
                                Y[b, m, h, w] += X[b, c, h + p, w + q] * W[m, c, p, q];  
                }  
    }
```

Computed  
by the grid

Computed by a thread



# Partial Kernel Code for a Convolution Layer

```
__global__ void ConvLayerForward_Basic_Kernel
(int C, int W_grid, int K, float* X, float* W, float* Y)
{
    int m = blockIdx.x;
    int h = (blockIdx.y / W_grid) * TILE_WIDTH + threadIdx.y;
    int w = (blockIdx.y % W_grid) * TILE_WIDTH + threadIdx.x;
    float acc = 0.0f;
    for (int c = 0; c < C; c++) {                // sum over all input channels
        for (int p = 0; p < K; p++)                // loop over KxK filter
            for (int q = 0; q < K; q++)
                acc += X[c, h + p, w + q] * W[m, c, p, q];
    }
    Y[m, h, w] = acc;
}
```

# Some Observations

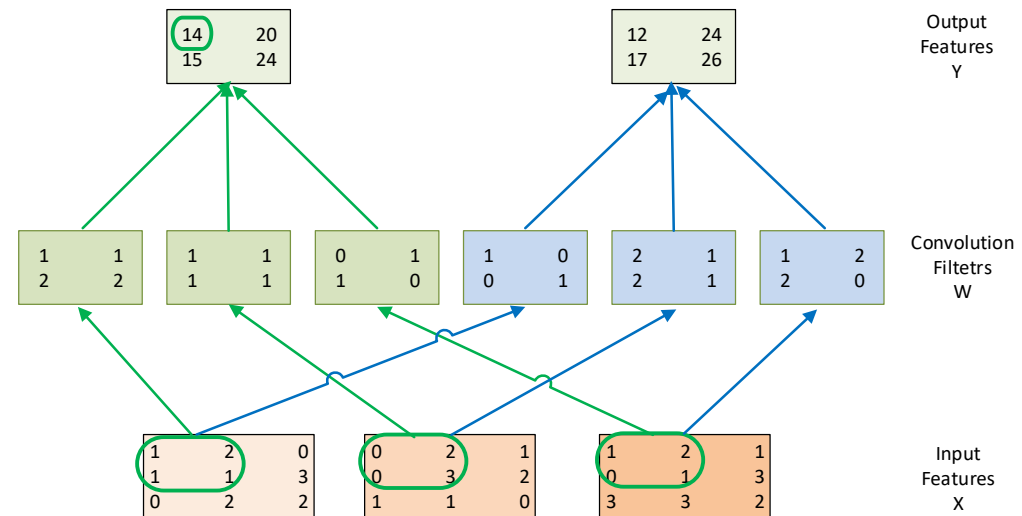
## Enough parallelism

- if the total number of pixels
- across all output feature maps is large
- (often the case for CNN layers)

## Each input tile

- loaded M times (number of output features), so
- **not efficient in global memory bandwidth,**
- but block scheduling in X dimension should give cache benefits.

# Implementing a Convolution Layer with Matrix Multiplication



1	1	2	2
1	0	0	1

1	1	1	1
2	1	2	1

0	1	1	0
1	2	2	0

Convolution Filters  $W'$

\*

1	2	1	1
2	0	1	3
1	1	0	2
1	3	2	2
0	2	0	3
2	1	3	2
0	3	1	1
3	2	1	0
1	2	1	1
2	1	0	3
0	1	3	3
1	3	3	2

Input Features  $X_{\text{unrolled}}$

=


14	20	15	24
12	24	17	26

Output Features  $Y$

# Simple Matrix Multiplication

Each product matrix element is an output feature map pixel.

This inner product generates element 0 of output feature map 0.

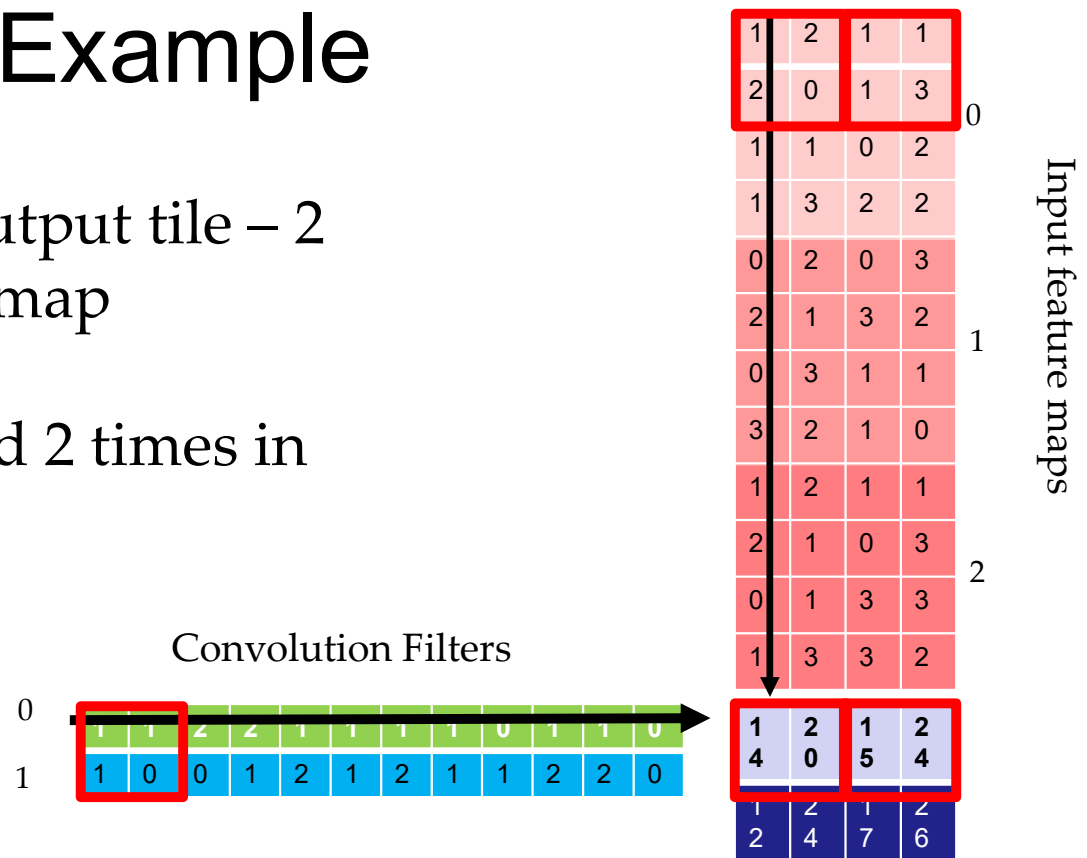


# Tiled Matrix Multiplication

## 2x2 Example

Each block calculates one output tile – 2 elements from each output map

Each input element is reused 2 times in the shared memory

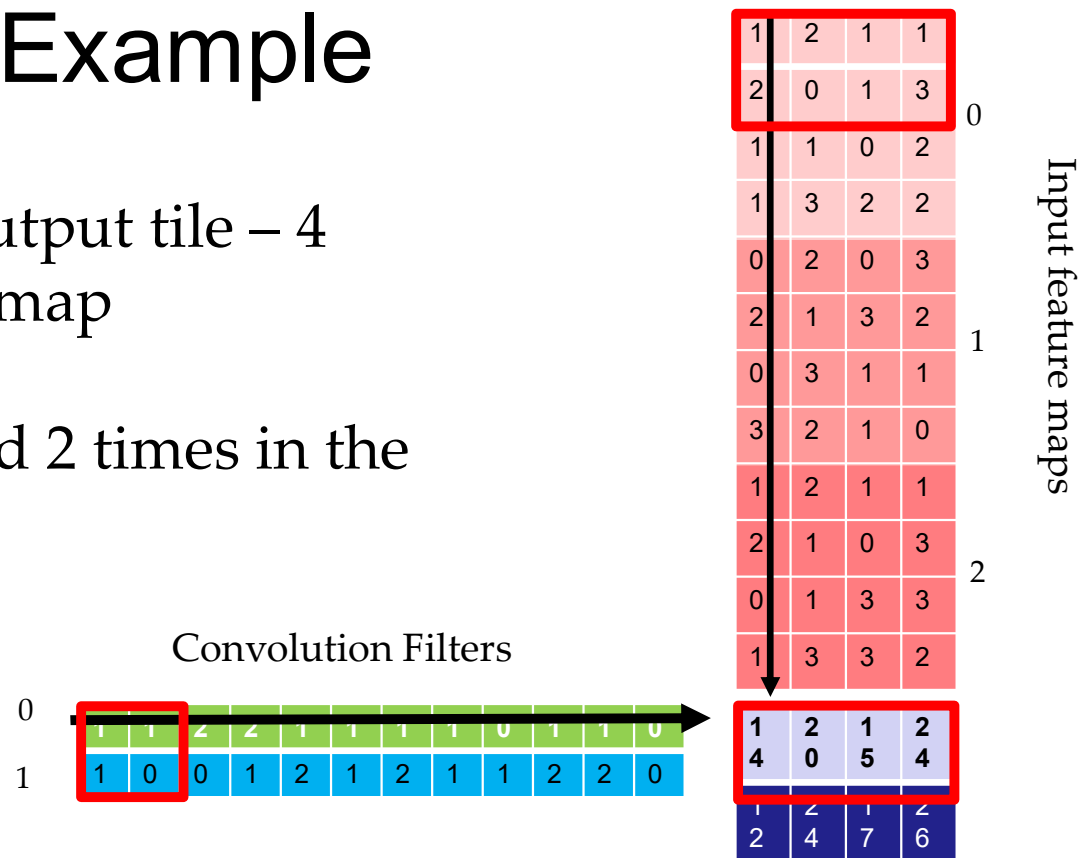


# Tiled Matrix Multiplication

## 2x4 Example


Each block calculates one output tile – 4 elements from each output map

Each input element is reused 2 times in the shared memory



# Efficiency Analysis: Total Input Replication

- Replicated input features are shared among output maps
  - There are  $H_{out} * W_{out}$  output feature map elements
  - Each requires  $K*K$  elements from the input feature maps
  - So, the total number of input element after replication is  $H_{out} * W_{out} * K * K$  times for each input feature map
  - The total number of elements in each original input feature map is  $(H_{in} + K - 1) * (W_{in} + K - 1)$



# Analysis of a Small Example

$$H_{\text{out}} = 2$$

$$W_{\text{out}} = 2$$

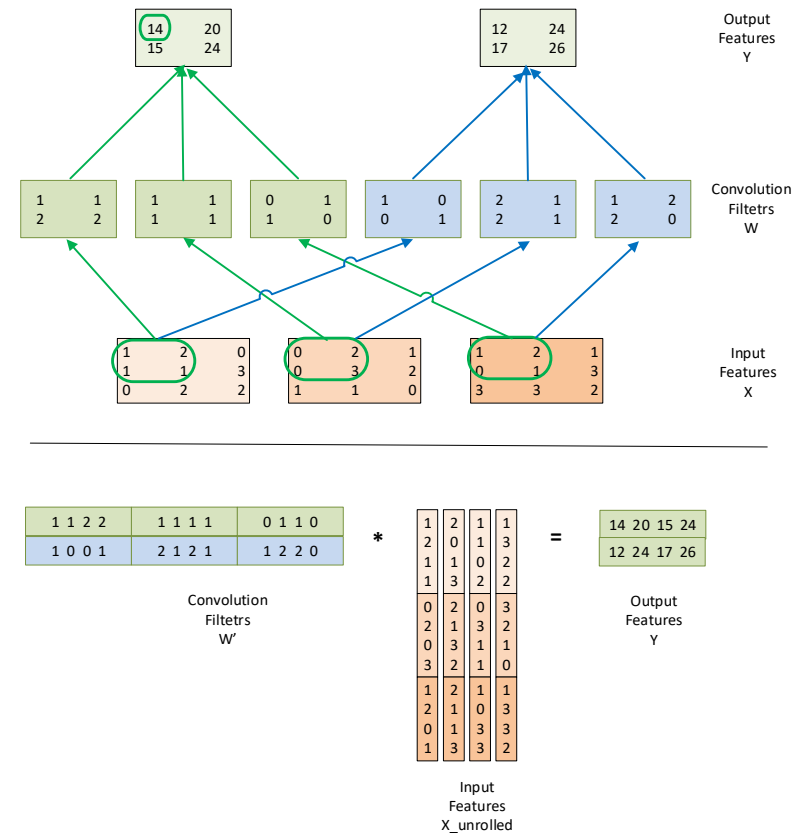
$$K = 2$$

There are 3 input maps (channels)

The total number of input elements in the replicated (“unrolled”) input matrix is  $3 \times 2 \times 2 \times 2 \times 2$

The replicating factor is

$$(3 \times 2 \times 2 \times 2 \times 2) / (3 \times 3 \times 3) = 1.78$$



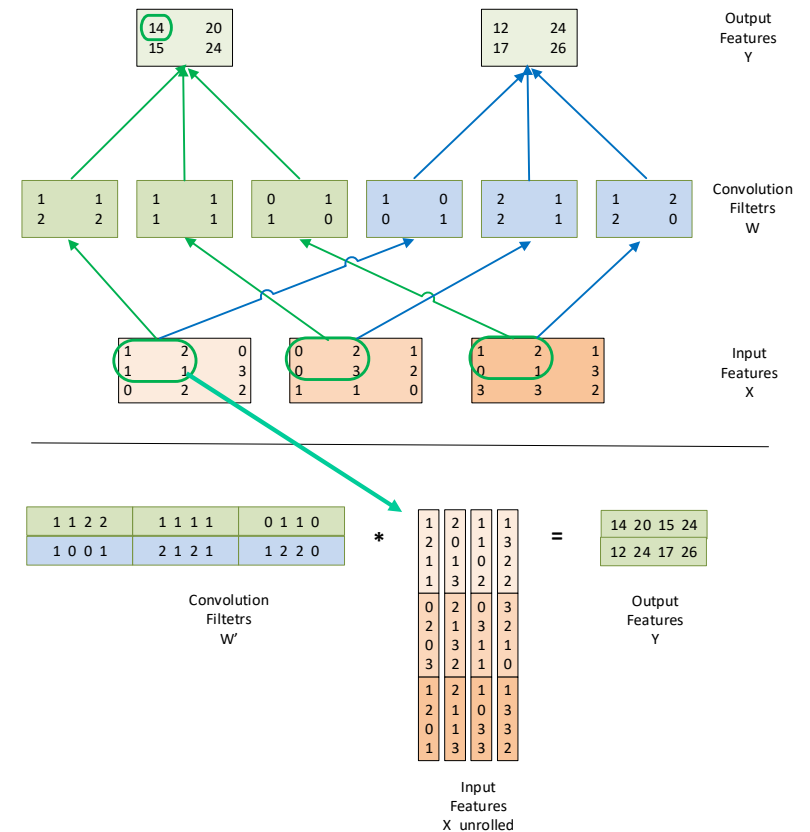


# Memory Access Efficiency of Original Convolution Algorithm

- Assume that we use tiled 2D convolution
- For input elements
  - Each output tile has  $\text{TILE\_WIDTH}^2$  elements
  - Each input tile has  $(\text{TILE\_WIDTH}+K-1)^2$
  - The total number of input feature map element accesses was  $\text{TILE\_WIDTH}^2 * K^2$
  - The reduction factor of the tiled algorithm is  $K^2 * \text{TILE\_WIDTH}^2 / (\text{TILE\_WIDTH}+K-1)^2$
- The convolution filter weight elements are reused within each output tile

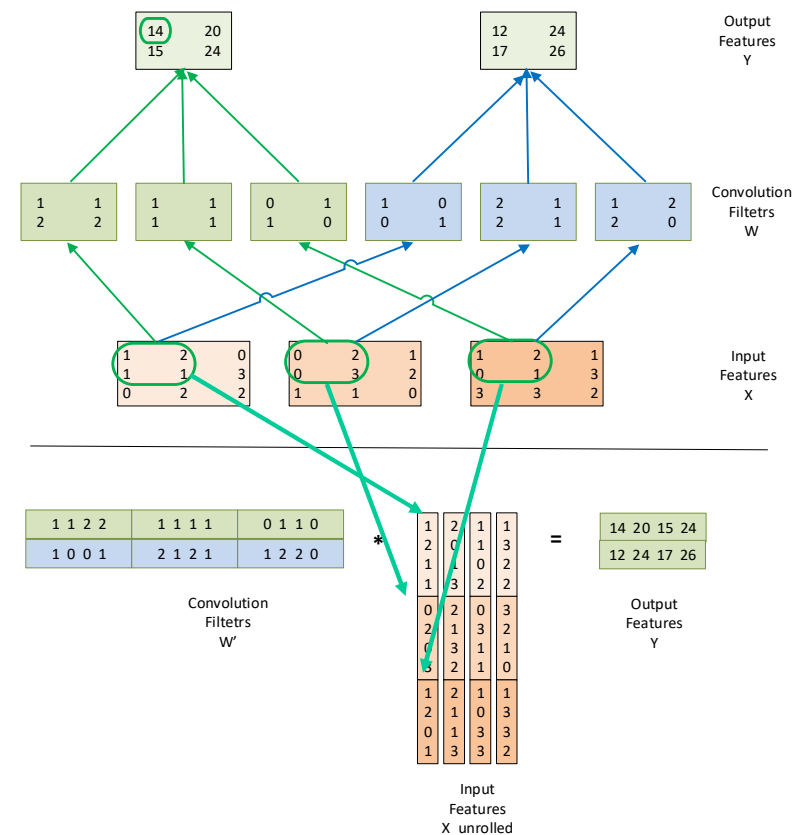
# Properties of the Unrolled Matrix

- Each unrolled column corresponds to an output feature map element
- For an output feature element  $(h,w)$ , the index for the unrolled column is  $h*W_{out}+w$  (linearized index of the output feature map element)



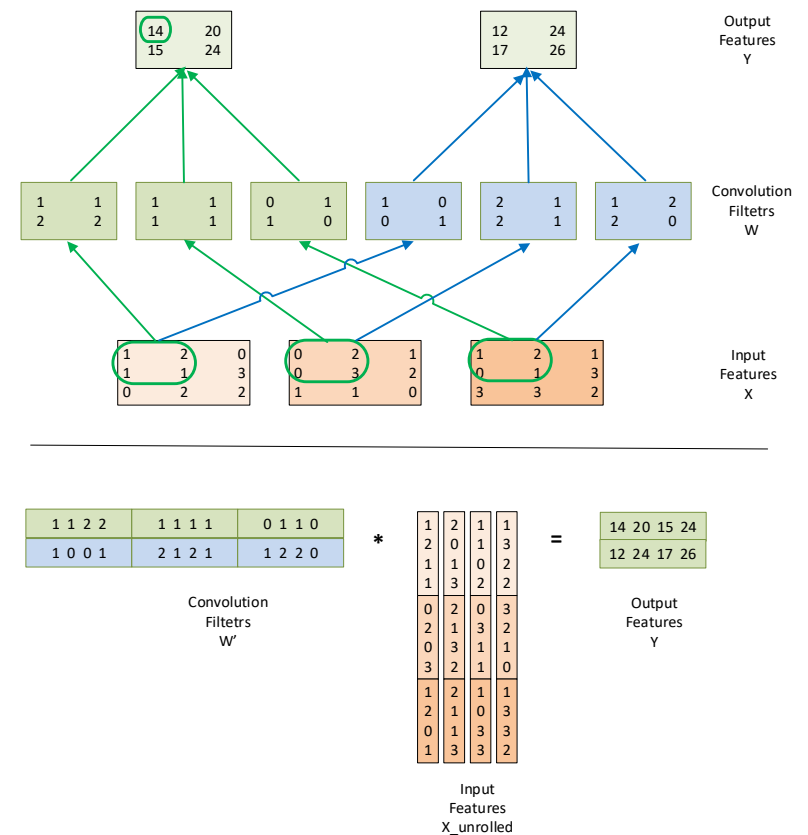
# Properties of the Unrolled Matrix (cont.)

- Each section of the unrolled column corresponds to an input feature map
- Each section of the unrolled column has  $k \times k$  elements (convolution mask size)
- For an input feature map  $c$ , the vertical index of its section in the unrolled column is  $c \times k \times k$  (linearized index of the output feature map element)



# To Find the Input Elements

- For output element  $(h, w)$ , the base index for the upper left corner of the input feature map  $c$  is  $(c, h, w)$
- The input element index for multiplication with the convolution mask element  $(p, q)$  is  $(c, h+p, w+q)$



# Input to Unrolled Matrix Mapping

Output element (h, w)

Mask element (p, q)

Input feature map c

```
// calculate the horizontal matrix index
```

```
int w_unroll = h * W_out + w;
```

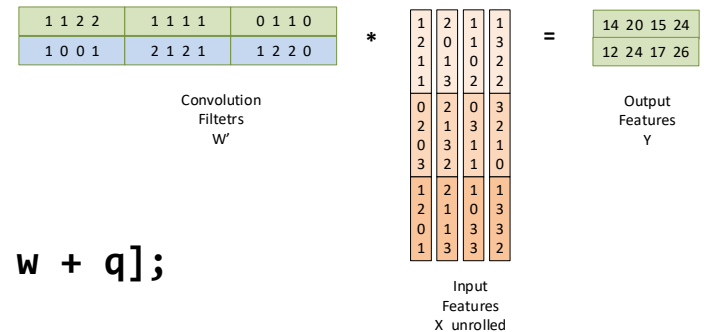

```
// find the beginning of the unrolled
```

```
int w_base = c * (K*K);
```

```
// calculate the vertical matrix index
```

```
int h_unroll = w_base + p * K + q;
```

```
X_unroll[b, h_unroll, w_unroll] = X[b, c, h + p, w + q];
```



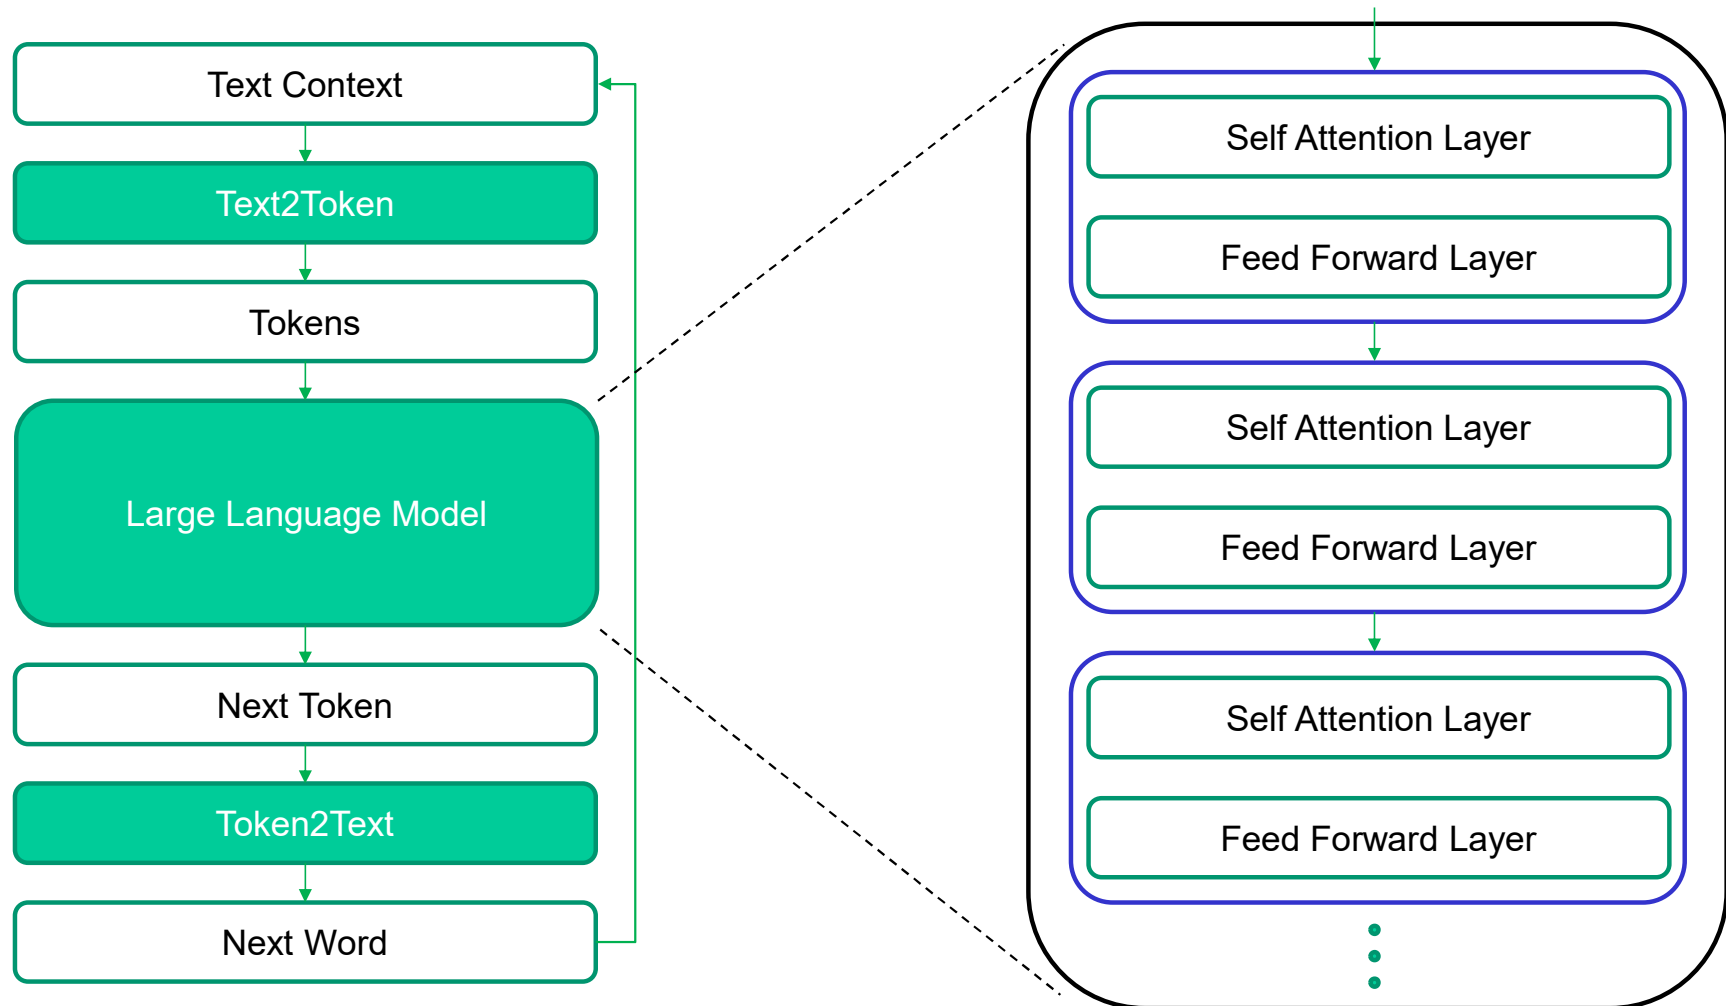
# Function to generate “unrolled” X

```
void unroll(int B, int C, int H, int W, int K, float* X, float* X_unroll)
{
    int H_out = H - K + 1;           // calculate H_out, W_out
    int W_out = W - K + 1;
    for (int b = 0; b < B; ++b)      // for each image
        for (int c = 0; c < C; ++c) { // for each input channel
            int w_base = c * (K*K);   // per-channel offset for smallest X_unroll index
            for (int p = 0; p < K; ++p) // for each element of KxK filter (two loops)
                for (int q = 0; q < K; ++q) {
                    for (int h = 0; h < H_out; ++h) // for each thread (each output value, two loops)
                        for (int w = 0; w < W_out; ++w) {
                            int h_unroll = w_base + p * K + q; // data needed by one thread
                            int w_unroll = h * W_out + w;        // smallest index--across threads (output values)
                            X_unroll[b, h_unroll, w_unroll] = X[b, c, h + p, w + q]; // copy input pixels
                        }
                }
        }
}
```

# Implementation Strategies for a Convolution Layer

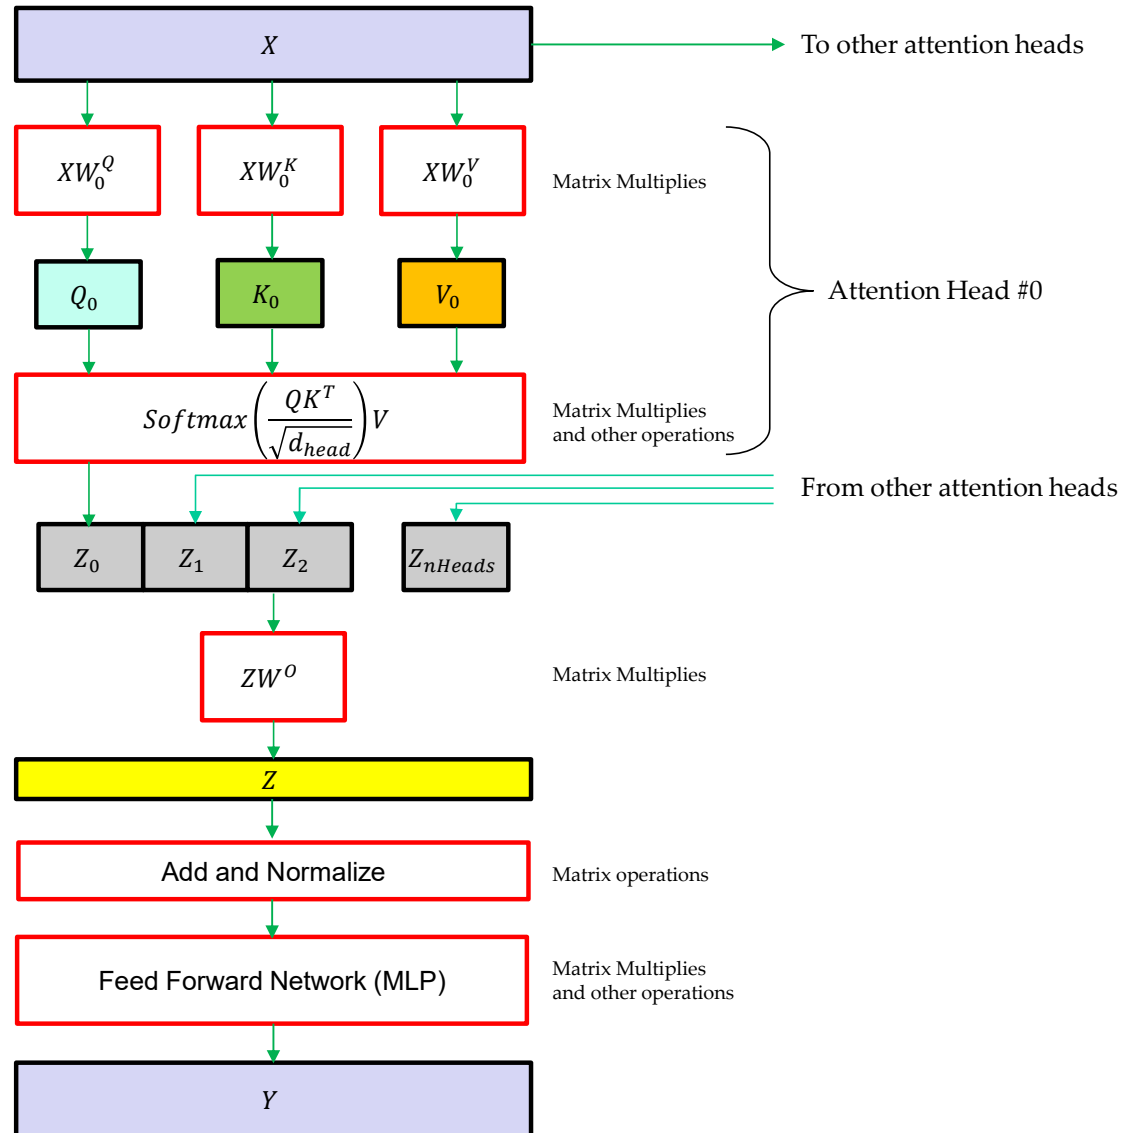

- **Baseline**
  - Tiled 2D convolution implementation, use constant memory for convolution masks
- **Matrix-Multiplication Baseline**
  - Input feature map unrolling kernel, constant memory for convolution masks as an optimization
  - Tiled matrix multiplication kernel
- **Matrix-Multiplication with built-in unrolling**
  - Perform unrolling only when loading a tile for matrix multiplication
  - The unrolled matrix is only conceptual
  - When loading a tile element of the conceptual unrolled matrix into the shared memory, use the properties in the lecture to load from the input feature map
- **More advanced Matrix-Multiplication**
  - Use joint register-shared memory tiling

# Transformer-based Language Models





# Single Layer Computational Flow




# GPT-3, as an example

Language Models are Few-Shot Learners, Brown et al., OpenAI, July 2020

Model Name	$n_{\text{params}}$	$n_{\text{layers}}$	$d_{\text{model}}$	$n_{\text{heads}}$	$d_{\text{head}}$	Batch Size	Learning Rate
GPT-3 Small	125M	12	768	12	64	0.5M	$6.0 \times 10^{-4}$
GPT-3 Medium	350M	24	1024	16	64	0.5M	$3.0 \times 10^{-4}$
GPT-3 Large	760M	24	1536	16	96	0.5M	$2.5 \times 10^{-4}$
GPT-3 XL	1.3B	24	2048	24	128	1M	$2.0 \times 10^{-4}$
GPT-3 2.7B	2.7B	32	2560	32	80	1M	$1.6 \times 10^{-4}$
GPT-3 6.7B	6.7B	32	4096	32	128	2M	$1.2 \times 10^{-4}$
GPT-3 13B	13.0B	40	5140	40	128	2M	$1.0 \times 10^{-4}$
GPT-3 175B or "GPT-3"	175.0B	96	12288	96	128	3.2M	$0.6 \times 10^{-4}$


**Table 2.1:** Sizes, architectures, and learning hyper-parameters (batch size in tokens and learning rate) of the models which we trained. All models were trained for a total of 300 billion tokens.

GPT-3 has 96 Layers, 55B parameters just from Self Attention. 220 Gflop per X vector, per output token.



In GPT-3, each  $W$  is 12288x128 matrix, 1.5M parameters, 3MB (@ 16 bit floats), or 6 MFlop per  $X$  vector

There are 96x3 of these for a total of 432M parameters, 864MB (@ 16 bit floats), or 1.7 GFlop per  $X$  vector




$W^O$  is 12288x12288 matrix, 150M parameters, 300MB (@ 16 bit floats), or 600 MFlop per  $X$  vector

# Project Overview

- Optimize the forward pass of the convolutional layers in a modified LeNet-5 CNN using CUDA. (CNN implemented using Mini-DNN, a C++ framework)
- The network will be classifying Fashion MNIST dataset
- Some network parameters to be aware of
  - Input Size: 86x86 pixels, batch of 1k-10k images
  - Input Channels: 1
  - Convolutional kernel size: 7x7
  - Number of kernels: Variable (your code should support this)



<https://github.com/zalandoresearch/fashion-mnist>



<http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>

A decorative graphic on the left side of the slide consisting of two vertical lines, one blue and one orange, extending from the top to the bottom of the slide.

# **QUESTIONS?**

# **READ CHAPTER 16!**

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483,  
University of Illinois, Urbana-Champaign